

Deployed hop-shell classifier (test sequential edits)

Metric	Value
Test n	2007
6-class accuracy	0.7005
Majority-class baseline	0.5291
Δ vs majority (pp)	17.14
Macro recall	0.4001
Macro F1	0.4049
Under-prediction rate	0.1908
Over-prediction rate	0.1086
Merged- $\{0,1\}$ (secondary)	0.7180
Selection	val_accuracy_guarded
Loss	ordinal $\lambda=0.5$